

```

Q1: Linear Regression      from sklearn.datasets import
fetch_california_housing from sklearn.model_selection import
train_test_split from sklearn.linear_model import
LinearRegression from sklearn.metrics import
mean_squared_error, r2_score import matplotlib.pyplot as plt
import pandas as pd

data = fetch_california_housing()
X = pd.DataFrame(data.data, columns=data.feature_names) y
= data.target
# (California dataset has no missing values, normalization optional)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)

model = LinearRegression() model.fit(X_train,
y_train)
print("Regression Coefficients:", model.coef_)
print("Intercept:", model.intercept_)

y_pred = model.predict(X_test) mse =
mean_squared_error(y_test, y_pred) r2 =
r2_score(y_test, y_pred) print(f"MSE:
{mse:.3f}") print(f"R2 Score: {r2:.3f}")

plt.scatter(y_test, y_pred, color='blue', alpha=0.5)
plt.xlabel("Actual Prices") plt.ylabel("Predicted
Prices") plt.title("Predicted vs Actual House
Prices") plt.show()

```

Q2: Logistic Regression

```

from sklearn.datasets import load_iris from
sklearn.model_selection import train_test_split from
sklearn.preprocessing import StandardScaler from
sklearn.linear_model import LogisticRegression from
sklearn.metrics import confusion_matrix, classification_report
import seaborn as sns import matplotlib.pyplot as plt

iris = load_iris()
X = iris.data y
= iris.target

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2,
random_state=42)

clf = LogisticRegression(max_iter=1000) clf.fit(X_train,
y_train)

print("Model Coefficients:\n", clf.coef_)

y_pred = clf.predict(X_test) print("\nConfusion Matrix:\n",
confusion_matrix(y_test, y_pred)) print("\nClassification Report:\n",
classification_report(y_test, y_pred))

plt.figure(figsize=(6,4))
sns.heatmap(confusion_matrix(y_test, y_pred), annot=True, fmt="d",
cmap="Blues") plt.xlabel("Predicted") plt.ylabel("Actual")
plt.title("Confusion Matrix - Iris Logistic Regression") plt.show()

```

### Q3: Decision Tree Classifier

```

from sklearn.tree import DecisionTreeClassifier,
plot_tree from sklearn.metrics import accuracy_score iris
= load_iris() X = iris.data y = iris.target

```

```

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2,
random_state=42)

clf = LogisticRegression(max_iter=1000) clf.fit(X_train,
y_train)

print("Model Coefficients:\n", clf.coef_)

for depth in [3, 5, None]: tree =
DecisionTreeClassifier(max_depth=depth, random_state=42)
tree.fit(X_train, y_train)

y_train_pred = tree.predict(X_train)
y_test_pred = tree.predict(X_test)

print(f"\nDecision Tree (max_depth={depth})") print("Training
Accuracy:", accuracy_score(y_train, y_train_pred)) print("Testing
Accuracy:", accuracy_score(y_test, y_test_pred))
plt.figure(figsize=(12,8)) plot_tree(tree,
feature_names=iris.feature_names,
class_names=iris.target_names, filled=True)
plt.show()

```

#### Q4: K-Means Clustering

```

from sklearn.cluster import KMeans from sklearn.preprocessing
import StandardScaler from sklearn.metrics import
silhouette_score, adjusted_rand_score import numpy as np

X = iris.data y_true
= iris.target

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

kmeans = KMeans(n_clusters=3, random_state=42) kmeans.fit(X_scaled)

```

```

print("Inertia:", kmeans.inertia_)

inertia = []
K = range(1, 10) for
k in K:
    km = KMeans(n_clusters=k, random_state=42).fit(X_scaled)
    inertia.append(km.inertia_)

plt.plot(K, inertia, marker='o')
plt.xlabel('k') plt.ylabel('Inertia')
plt.title('Elbow Method')
plt.show()

labels = kmeans.labels_ ari =
adjusted_rand_score(y_true, labels) silhouette
= silhouette_score(X_scaled, labels)
print(f"Adjusted Rand Index: {ari:.3f}")
print(f"Silhouette Score: {silhouette:.3f}")

plt.scatter(X_scaled[:, 0], X_scaled[:, 1], c=labels, cmap='viridis', alpha=0.6)
plt.scatter(kmeans.cluster_centers_[:, 0], kmeans.cluster_centers_[:, 1],
color='red', marker='X', s=200, label='Centroids') plt.title("K-Means
Clustering on Iris") plt.legend() plt.show()

```

## Q5. Cross Validation & Ensemble

```

from sklearn.datasets import load_breast_cancer
from sklearn.svm import SVC from sklearn.ensemble import
RandomForestClassifier from sklearn.model_selection import
cross_val_score, train_test_split from sklearn.metrics import
confusion_matrix, classification_report import numpy as np

data = load_breast_cancer()
X = data.data y =
data.target

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)

```

```

svm = SVC(kernel='rbf', gamma='scale') scores =
cross_val_score(svm, X_train, y_train, cv=5) print("SVM
Cross-Validation Accuracies:", scores) print(f"Average:
{scores.mean():.3f}, Std: {scores.std():.3f}")
t
svm.fit(X_train, y_train) print("SVM Test Accuracy:",
svm.score(X_test, y_test))

rf = RandomForestClassifier(n_estimators=100, random_state=42)
rf.fit(X_train, y_train) y_pred = rf.predict(X_test)

print("\nRandom Forest Test Accuracy:", rf.score(X_test, y_test))
print("\nConfusion Matrix:\n", confusion_matrix(y_test, y_pred))
print("\nClassification Report:\n", classification_report(y_test, y_pred))

rf_scores = cross_val_score(rf, X_train, y_train, cv=5) print("\nRandom
Forest Cross-Validation Accuracy:", rf_scores) print(f"Average:
{rf_scores.mean():.3f}, Std: {rf_scores.std():.3f}")

importances = rf.feature_importances_
indices = np.argsort(importances)[-10:]

plt.figure(figsize=(10,6)) plt.bar(range(10),
importances[indices[:10]], align="center")
plt.xticks(range(10), [data.feature_names[i] for i in
indices[:10]], rotation=45) plt.title("Top 10 Feature
Importances (Random Forest)") plt.show()

```